EXPERIMENTAL METHOD

One baseline. One change. Three independent checks.

A frozen BF16 decode workload; the only optimized mechanism is the Small-M GEMM route.

01

FIX INPUTS

Fixed SGLang revision
Model + request set

02

IDENTIFY COST

Nsight Systems trace
Native GEMM geometry

03

ADAPT ROUTE

Opt-in BF16 Small-M
Preserve FP32 bias

04

VERIFY + MEASURE

Correctness and throughput
Interleaved A/B pairs

VALIDATION

Numerical identity

340 GEMM checks
255 batch / reorder checks

Executed path

Nsight kernel launches
Native vs candidate geometry

End-to-end result

5 paired server runs
Median throughput +72.4%